

# Architecture & Design Principles

## Overview


The coordinator-node-starter is a template system for building competition coordinators ("Crunches"). Each coordinator runs a predict → score → rank loop against participant models, fed by live market data.

The architecture is built on five core principles:


1. Canonical schema with JSONB extension
2. Deterministic callable wiring via packs
3. Generated skill files that guide LLM agents
4. Shared runtime library ( `coordinator_runtime` )
5. Two-repo separation per competition

## 1. Canonical Schema + JSONB Extension

Every coordinator shares the same Postgres table structure. Protocol-required columns are typed and indexed. Competition-specific data lives in JSONB columns.



The key insight: **typed columns for protocol queries, JSONB columns for competition-specific payloads.**



|   |                                     |  |  |
|---|-------------------------------------|--|--|
| scope_jsonb:<br>{asset,<br>horizon}                       | scope_jsonb:<br>{corpus,<br>window} | scope_jsonb:<br>{universe,<br>rebalance} |  |
| Zero schema migrations. Same queries. Different payloads. |                                     |  |  |

Separately, `market_records` is a **high-volume append table** for raw market data. Each data point is its own row with typed/indexed columns for querying (`provider`, `asset`, `kind`, `granularity`, `ts_event`). JSONB is only used for the per-record value shape, which varies by data kind:

| MARKET_RECORDS |              |    |                          |
|----------------|--------------|----|--------------------------|
| string         | id           | PK |                          |
| string         | provider     |    | indexed                  |
| string         | asset        |    | indexed                  |
| string         | kind         |    | indexed (tick or candle) |
| string         | granularity  |    | indexed                  |
| datetime       | ts_event     |    | indexed                  |
| datetime       | ts_ingested  |    | indexed                  |
| jsonb          | values_jsonb |    | ← varies by kind         |
| jsonb          | meta_jsonb   |    |                          |



| MARKET_INGESTION_STATE |               |    |           |
|------------------------|---------------|----|-----------|
| string                 | id            | PK |           |
| string                 | provider      |    | indexed   |
| string                 | asset         |    | indexed   |
| string                 | kind          |    | indexed   |
| string                 | granularity   |    | indexed   |
| datetime               | last_event_ts |    | watermark |

```

tick row:  values_jsonb = {"price": 44987.5}
candle row: values_jsonb = {"open": 45000, "high": 45010, "low": 44980, "close": 44992, "volume": 12.3}

```

This is not a JSONB-array-per-feed design. One row per data point, uniquely indexed on `(provider, asset, kind, granularity, ts_event)` for dedup and fast time-range queries.

This means:

- **New competitions don't require schema migrations.** A BTC price predictor and an NLP sentiment ranker use the same tables — they just put different shapes inside the JSONB columns.
- **Core queries work everywhere.** Filtering predictions by model, time range, or scope works identically regardless of what's inside the payloads.
- **The webapp, CLI tools, and report APIs work against known column names** while competition-specific UI (column labels, metric widgets) reads from configurable schema endpoints.

The canonical tables are defined once in `coordinator_core/infrastructure/db/db_tables.py` and shared across all generated workspaces via vendoring.


## 2. Deterministic Callable Wiring via Packs

A coordinator's behavior is defined by **ten callable slots** that form the full predict → score → rank pipeline:



| Slot                       | Responsibility   |
|----------------------------|--|
| RAW_INPUT_PROVIDER         | Fetch/build the raw data payload for a prediction cycle              |
| INFERENCE_INPUT_BUILDER    | Transform raw input into what models receive                         |
| INFERENCE_OUTPUT_VALIDATOR | Validate model responses   |
| SCORING_FUNCTION           | Score a single prediction against ground truth                       |
| GROUND_TRUTH_RESOLVER      | Resolve what actually happened (from market tape or external source) |
| MODEL_SCORE_AGGREGATOR     | Aggregate per-prediction scores into per-model scores                |
| LEADERBOARD_RANKER         | Rank models into a leaderboard snapshot                              |
| REPORT_SCHEMA_PROVIDER     | Define the leaderboard column / metrics widget schema for the UI     |
| PREDICTION_SCOPE_BUILDER   | Build the scope dimensions for each prediction                       |
| PREDICT_CALL_BUILDER       | Assemble the final model invocation payload                          |

**Packs** are JSON manifests that pre-wire all ten slots:




Available packs:

- **baseline** — 60s cycles, balanced local development
- **realtime** — 15s cycles, low-latency tournament
- **tournament** — combined in-sample + out-of-sample scopes

Callable paths use Python dotted-module notation (`runtime_definitions.data:provide_raw_input`). Workers load them at startup via `importlib`, so no code changes are needed to swap behavior — just change the env var.

### 3. Generated Skill Files (LLM Agent Guidance)

Every scaffolded workspace includes three `SKILL.md` files — structured instructions for coding agents:



The idea: `coordinator init` produces the hard 80% (working infrastructure, wired callables, running stack). The SKILL files prompt an LLM to fill in the remaining 20% — the actual competition logic — by pointing it at the right files and explaining the constraints.

---


## 4. Shared Runtime Library ( `coordinator_runtime` )

Code that is **common across all coordinators** lives in `coordinator_runtime/`:




### Feed Architecture


Provider adapters implement a dual-mode protocol:



Provider adapters prefer SDK packages (e.g. `python-binance`) for external connections, with HTTP fallback if the package is unavailable.


## Tape-First Pattern





## 5. Two-Repo Separation

Each competition produces two repositories with a clear visibility boundary:



| Repo                                  | Visibility     | Contains   |
|---------------------------------------|----------------|--|
| <code>crunch-&lt;name&gt;</code>      | <b>Public</b>  | Model interface ( <code>tracker.py</code> ), scoring logic, quickstarter examples, schemas |
| <code>crunch-node-&lt;name&gt;</code> | <b>Private</b> | Deployment config, callable wiring, runtime definitions, vendored runtime packages         |

This separation means participants see only the challenge contract — not the infrastructure, scoring weights, or data pipeline internals.

The node workspace is self-contained: it vendors `coordinator_core`, `coordinator_runtime`, and `node_template` into a `runtime/` directory and builds a single Docker image from them. No dependency on the template repo at runtime.


---

## How It Fits Together

### Scaffold Output

```
coordinator init btc-trader --pack realtime
|
└ reads pack JSON (callables, schedule, checkpoint interval)
└ renders templates (Dockerfile, docker-compose, Makefile, SKILL files)
└ generates .local.env + config/callables.env with all 10 slots wired
└ vendors coordinator_core + coordinator_runtime + node_template
|
└ btc-trader/
    └ SKILL.md                                ← workspace agent guidance
    └ process-log.jsonl                         ← scaffold audit trail
    |
    └ crunch-btc-trader/                      ← public challenge package
        └ SKILL.md                            ← challenge agent guidance
        └ crunch_btc_trader/
            └ tracker.py                     ← model interface (fill this in)
            └ scoring.py                      ← scoring function (fill this in)
            └ examples/                        ← quickstarters for participants
            └ pyproject.toml
    |
    └ crunch-node-btc-trader/                 ← private node runtime
        └ SKILL.md                            ← node agent guidance
        └ RUNBOOK.md                          ← troubleshooting guide
        └ docker-compose.yml                  ← all services pre-wired
        └ Dockerfile                           ← deploy/verify/logs targets
        └ Makefile                            ← environment config
        └ .local.env                           ← node-private callables
        └ runtime_definitions/
            └ data.py                            ← tape-first input + ground truth
            └ inference.py                     ← input builder
            └ validation.py                    ← output validator
            └ reporting.py                     ← UI schema
        └ config/
            └ callables.env                     ← all 10 callable paths
            └ scheduled_prediction_configs.json
        └ runtime/                             ← vendored packages
            └ coordinator_core/
            └ coordinator_runtime/
            └ node_template/
```

### Running Stack



## Key Design Decisions

| Decision                                | Rationale   |
|---|---|
| JSONB for all competition-specific data | No migrations per competition. Same queries, different payloads.              |
| Ten callable slots, not plugin classes  | Simple to reason about. One function per concern. Swappable via env.          |
| Packs, not inheritance                  | Flat JSON config. No class hierarchies. Easy to diff and override.            |
| Tape-first ground truth                 | Deterministic scoring from stored data. Feed failures don't block scoring.    |
| Vendor at scaffold time                 | Node workspace works offline. No implicit dependency on template repo.        |
| Generated SKILL files                   | Agents get structured guidance, not just code. The scaffold is documentation. |
| SDK-first for external APIs             | Explicit, testable dependencies. HTTP fallback for resilience.                |
| Feed visibility in UI is data-driven    | Feeds nav appears only when feed data exists. Zero config for the operator.   |